

Sid Narayan

SOFTWARE ENGINEER

San Francisco, CA

☎ (+1) 626-230-0465 | ✉ sidsdn@protonmail.com | 📷 sidheart | 🌐 sidheartn

Education

UCLA (University of California, Los Angeles)

B.S. IN COMPUTER SCIENCE

- GPA: 3.6 /4.0

Los Angeles, California

Sep. 2013 - Dec. 2017 (Expected)

Experience

Amazon

SOFTWARE DEVELOPMENT ENGINEER 2 (AMAZON ELASTICSEARCH)

Palo Alto, CA

Mar. 2020 - Jan. 2020

- Designed and implemented a self-healing feature for Elasticsearch using Java which detects node and cluster issues and fixes them in real time.
- Developed a dynamic throttling mechanism for Elasticsearch which stops floods of queries and requests from taking down any cluster.
- Contributed to open source Elasticsearch plugins for metrics collection, aggregation, and analytics.

Splunk

SOFTWARE ENGINEER (PLATFORM TEAM)

San Jose, CA

Oct. 2019 - Mar. 2020

- Improved the routing logic in our Golang based control plane, giving customers fine-grained control over where their data resides and who can access it.
- Implemented a new authentication service backed by a distributed cache and Kafka, removing a critical system bottleneck and reducing average user authentication time by 75%.
- Refactored a Java based REST API layer into distributed microservices in order to allow our system to scale horizontally.

SignalFx

SOFTWARE ENGINEER (METADATA/APPLICATION PLATFORM TEAMS)

San Mateo, CA

May. 2018 - Oct. 2019

- Shipped a distributed service backed by Helix and Elasticsearch which matches event streams with user-defined rules in order to mute annoying alerts in real time.
- Integrated MongoDB with SignalFx's backend architecture and wrote a generic Java client to extract, transform, and load business objects into the cluster.
- Built a cloud storage migration system (EBS, GCS, etc.) which reduced our Elasticsearch, Cassandra, and MongoDB cluster average MTTR by 90%.
- Integrated Facebook's zstandard compression library with our backup pipeline, dramatically reducing file size and saving \$3000/month.

Illumio

SOFTWARE ENGINEER (PRODUCTION ENGINEERING INTERN)

Sunnyvale, CA

Jun. 2016 - Sept. 2016

- Leveraged Ruby on Rails, Scala, and Gatling to create a system for managing, monitoring, analyzing, and executing all production tests.
- Built a REST API server to allow production engineers to build and execute complex testing scripts.
- Integrated the system with the company's OneLogin authentication to ensure that only production engineers could use the system.

Intel

SOFTWARE ENGINEER (NETWORK PLATFORM LAB INTERN)

Hillsboro, OR

Jun. 2015 - Oct. 2015

- Utilized Python and Django to quickly collect and analyze gigabytes of server data logs and display the statistics on a web server.
- Implemented and modified core control plane components of a 4G LTE cellular network using C++.
- Discovered and optimized inefficient writes in the software, resulting in an application speedup of 200%.

Awesome Projects

Expedia Big Data Analysis

PYTHON, PANDAS, NUMPY, D3.JS

- Analyzed a LARGE (multiple TB) Expedia dataset using Python and used the results to develop a marketing strategy to improve customer conversion.
- Developed a visualization to demonstrate and showcase the data to a panel of judges using Javascript and D3.
- Finalist project at UCLA's DataFest 2017!

Connoisseur (Capstone)

IOS, REACT NATIVE, MOBILE

- Built an iOS application using the React Native framework which provides restaurant suggestions based on a Netflix-like recommendation algorithm.
- Migrated from Node.js to a Java Spring backend infrastructure, resulting in more efficient API calls.
- Overhauled the front-end interface to provide a robust visual design and improve usability and functionality.